SOCIAL DATA SCIENCE

Tidy Data, Data Manipulation & Functions

Sebastian Barfort

August 06, 2016

University of Copenhagen
Department of Economics

"Herein lies the dirty secret about most data scientists' work – it's more data munging than deep learning. The best minds of my generation are deleting commas from log files, and that makes me sad. A Ph.D. is a terrible thing to waste."

Source

TECHNOLOGY

For Big-Data Scientists, 'Janitor Work' Is Key Hurdle to Insights

By STEVE LOHR   AUG. 17, 2014

Technology revolutions come in measured, sometimes foot-dragging steps. The lab science and marketing enthusiasm tend to underestimate the bottlenecks to progress that must be overcome with hard work and practical engineering.

The field known as "big data" offers a contemporary case study. The catchphrase stands for the modern abundance of digital data from many sources — the web, sensors, smartphones and corporate databases — that can be mined with clever software for discoveries and insights. Its promise is smarter, data-driven decision-making in every field. That is why data scientist is the economy's hot new job.

Yet far too much handcrafted work — what data scientists call "data wrangling," "data

Monica Rogati, Jawbone's vice president for data science, with Brian Wilt, a senior data scientist. Peter DaSilva for The New York Times

Source

RAW VERSUS PROCESSED DATA

Raw data

The original source of the data

Often hard to use directly for data analysis

You should *never* process your original data

Processed data

Data that is ready for analysis

Data manipulation involves going from *raw* to *processed* data.

This can include merging, subsetting, transforming, etc.

*All* steps that take you from raw to processed data should be scripted

Introduce some tricks for working (efficiently) with data

Introduce concept and tools for working with tidy data (`tidyr`)

Manipulate tidy data using `dplyr`

Iterate over elements using functions (`purrr`)

String processing (`stringr`, regular expressions)

The Pipe

magrittr::

%>%

The pipe operator %>% (RStudio has keyboard shortcuts, learn to use them!) let's you write sequences instead of nested functions

x %>% f(y) -> f(x,y)

x %>% f(z, .) -> f(z, x)

Read %>% as "then". First do this, *then* do this, etc...

It's implemented in R by a Danish econometrician

All the packages you will learn today work with the pipe.

```
enjoy(cool(bake(shape(beat(append(bowl(rep("flour", 2),
"yeast", "water", "milk", "oil"), "flour", until =
"soft"), duration = "3mins"), as = "balls", style =
"slightly-flat"), degrees = 200, duration = "15mins"),
duration = "5mins"))
```

```
bowl(rep("flour", 2), "yeast", "water", "milk", "oil") %>%
    append("flour", until = "soft") %>%
    beat(duration = "3mins") %>%
    shape(as = "balls", style = "slightly-flat") %>%
    bake(degrees = 200, duration = "15mins") %>%
    cool(buns, duration = "5mins") %>%
    enjoy()
```

source

Tidy data

tidyr

Tidy data: observations are in the rows, variables are in the columns

`tidyr`: take your messy data and turn it into a tidy format

Advantages of tidy data:

- Consistency
- Allows you to spend more time on your analysis
- Speed

variables   observations   values

- `gather`: Reshape from wide to long
- `spread`: Reshape from long to wide
- `separate`: Split a variable into multiple variables.

(Also more complicated functions such as `nest` for nested data frames, but we won't go into detail with those here)

```r
library("readr")
gh.link = "https://raw.githubusercontent.com/"
user.repo = "hadley/tidyr/"
branch = "master/"
link = "vignettes/pew.csv"
data.link = paste0(gh.link, user.repo, branch, link)
df = read_csv(data.link)
```

First five columns

religion	<$10k	$10-20k	$20-30k	$30-40k
Agnostic	27	34	60	81
Atheist	12	27	37	52
Buddhist	27	21	30	34

Question 1: What variables are in this dataset?

Question 2: How does a tidy version of this data look like?

Problem: Column names are not names of a variable, but *values* of a variable.

Objective: Reshaping wide format to long format

To tidy such data, we need to **gather** the non-variable columns into a two-column key-value pair

gather

Three parameters

1. Set of columns that represent values, not variables.
2. Name of the variable whose values form the column names (key).
3. The name of the variable whose values are spread over the cells (value.

```
library("tidyr")
args(gather)

## function (data, key, value, ..., na.rm = FALSE, conve
##     factor_key = FALSE)
## NULL
```

```
df.gather = df %>%
  gather(key = income,
         value = frequency,
         -religion)
```

religion	income	frequency
Agnostic	<$10k	27
Atheist	<$10k	12
Buddhist	<$10k	27
Catholic	<$10k	418
Don't know/refused	<$10k	15

ALTERNATIVES

This

```
df %>%
  gather(key = income,
         value = frequency,
         2:11)
```

returns the same as

```
df %>%
  gather(key = income,
         value = frequency,
         -religion)
```

Billboard data

```
library("readr")
gh.link = "https://raw.githubusercontent.com/"
user.repo = "hadley/tidyr/"
branch = "master/"
link = "vignettes/billboard.csv"
data.link = paste0(gh.link, user.repo, branch, link)
df = read_csv(data.link)
```

```
df[1:5, 1:5]
```

year	artist	track	time	date.entered
2000	2 Pac	Baby Don't Cry (Keep...	4:22	2000-02-26
2000	2Ge+her	The Hardest Part Of ...	3:15	2000-09-02
2000	3 Doors Down	Kryptonite	3:53	2000-04-08
2000	3 Doors Down	Loser	4:24	2000-10-21
2000	504 Boyz	Wobble Wobble	3:35	2000-04-15

```
df[1:5, 6:10]
```

wk1	wk2	wk3	wk4	wk5
87	82	72	77	87
91	87	92	NA	NA
81	70	68	67	66
76	76	72	69	67
57	34	25	17	17

Question: what are the variables here?

To tidy this dataset, we first gather together all the wk columns. The column names give the week and the values are the ranks:

```
billboard2 = df %>%
  gather(key = week,
         value = rank, wk1:wk76,
         na.rm = TRUE)
```

Not displaying the `track` column

year	artist	time	date.entered	week	rank
2000	2 Pac	4:22	2000-02-26	wk1	87
2000	2Ge+her	3:15	2000-09-02	wk1	91
2000	3 Doors Down	3:53	2000-04-08	wk1	81
2000	3 Doors Down	4:24	2000-10-21	wk1	76
2000	504 Boyz	3:35	2000-04-15	wk1	57

Are we done?

DATA CLEANING

Let's turn the week into a numeric variable and create a proper date column

```
library("dplyr")
billboard3 = billboard2 %>%
  mutate(
    week = extract_numeric(week),
    date = as.Date(date.entered) + 7 * (week - 1)) %>%
  select(-date.entered) %>%
  arrange(artist, track, week)
```

What functions from tidyr did we use here?

year	artist	track	time	week
2000	2 Pac	Baby Don't Cry (Keep…	4:22	1
2000	2 Pac	Baby Don't Cry (Keep…	4:22	2
2000	2 Pac	Baby Don't Cry (Keep…	4:22	3
2000	2 Pac	Baby Don't Cry (Keep…	4:22	4
2000	2 Pac	Baby Don't Cry (Keep…	4:22	5

After gathering columns, the key column is sometimes a combination of multiple underlying variable names.

```
library("readr")
gh.link = "https://raw.githubusercontent.com/"
user.repo = "hadley/tidyr/"
branch = "master/"
link = "vignettes/tb.csv"
data.link = paste0(gh.link, user.repo, branch, link)
df = read_csv(data.link)
```

iso2	year	m04	m514	m014	m1524	m2534	m3544
AD	1989	NA	NA	NA	NA	NA	NA
AD	1990	NA	NA	NA	NA	NA	NA
AD	1991	NA	NA	NA	NA	NA	NA
AD	1992	NA	NA	NA	NA	NA	NA
AD	1993	NA	NA	NA	NA	NA	NA

Question: what are the variables here?

The dataset comes from the World Health Organisation, and records the counts of confirmed tuberculosis cases by country, year, and demographic group. The demographic groups are broken down by sex (m, f) and age (0-14, 15-25, 25-34, 35-44, 45-54, 55-64, unknown).

```
tb2 = df %>%
  gather(demo, n, -iso2, -year, na.rm = TRUE)
```

iso2	year	demo	n
AD	2005	m04	0
AD	2006	m04	0
AD	2008	m04	0
AE	2006	m04	0
AE	2007	m04	0

Is this dataset tidy?

separate makes it easy to split a variable into multiple variables.
You can either pass it a regular expression to split on or a vector of
character positions. In this case we want to split after the first
character.

```
tb3 = tb2 %>%
  separate(demo, c("sex", "age"), 1)
```

iso2	year	sex	age	n
AD	2005	m	04	0
AD	2006	m	04	0
AD	2008	m	04	0
AE	2006	m	04	0
AE	2007	m	04	0

There are times when we are required to turn long formatted data into wide formatted data. The `spread` function spreads a key-value pair across multiple columns.
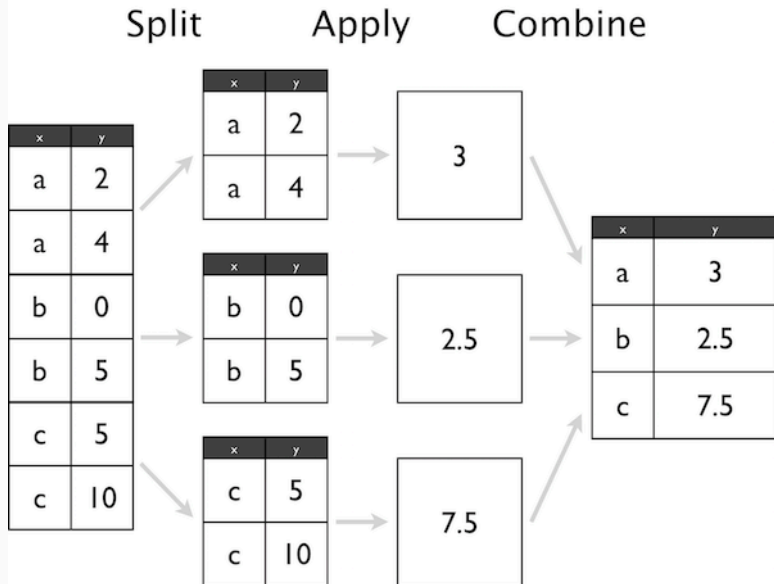
```
tb3.wide = tb3 %>% spread(age, n)
```

iso2	year	sex	014	04	1524	2534	3544
AD	1996	f	0	NA	1	1	0
AD	1996	m	0	NA	0	0	4
AD	1997	f	0	NA	1	2	3
AD	1997	m	0	NA	0	1	2
AD	1998	m	0	NA	0	0	1

Data Manipulation

Once you have your data stored in tidy form, you can easily apply a
split-apply-combine strategy, where you break up a big problem into
manageable pieces, opereate on each piece independently and then
put the pieces back together

Journal of Statistical Software

April 2011, Volume 40, Issue 1.   http://www.jstatsoft.org/

The Split-Apply-Combine Strategy for Data
Analysis

Hadley Wickham
Rice University

dplyr: (efficiently) split-apply-combine for data frames

Verbs a verb is a function that takes a data frame as it's first argument

- · filter: select rows
- · arrange: order rows
- · select: select columns
- · rename: rename columns
- · distinct: find distinct rows
- · mutate: add new variables
- · summarise: summarize across a data set
- · sample_n: sample from a data set

DATA

In this part of the lecture we will work with the Danish federal budget proposal for 2016

```
library("readr")
library("dplyr")
gh.link = "https://raw.githubusercontent.com/"
user.repo = "sebastianbarfort/sds_summer/"
branch = "gh-pages/"
link = "data/finanslov_tidy.csv"
data.link = paste0(gh.link, user.repo, branch, link)
df = read_csv(data.link)
```

Some nice guy has already cleaned this data for you

Try yourself

```
View(df)
glimpse(df)
summary(df)
head(df)
```

filter return rows with matching conditions.

```
df.max.udgift = df %>%
  filter(udgift == max(udgift)) %>%
  select(paragraf, aar, udgift)
```

paragraf	aar	udgift
Beskæftigelsesministeriet	2018	132541.4

```
df.skat = df %>%
  filter(paragraf == "Skatter og afgifter") %>%
  select(paragraf, aar, udgift) %>%
  arrange(-udgift)
```

paragraf	aar	udgift
Skatter og afgifter	2014	14487.1
Skatter og afgifter	2015	14401.6
Skatter og afgifter	2016	14386.2
Skatter og afgifter	2014	185.9
Skatter og afgifter	2015	185.9

CREATING NEW VARIABLES

mutate let's you add new variables to your data frame

```
df.mutated = df %>%
  mutate(newVar = udgift/2) %>%
  select(newVar, udgift)
```

newVar	udgift
38.85	77.7
13.20	26.4
193.90	387.8
132.10	264.2
3.25	6.5

SAMPLING FROM A DATA FRAME

We can sample from a data frame using sample_n and
sample_frac

```
df.sample.n = df %>%
  select(paragraf, aar, udgift) %>%
  sample_n(3)
```

paragraf	aar	udgift
Erhvervs- og Vækstministeriet	2017	-26.0
Transport- og Bygningsministeriet	2016	5.4
Beskæftigelsesministeriet	2016	0.6

GROUPED OPERATIONS

So far, we have primarily learned how to manipulate data frames.

The dplyr package becomes really powerful when we introduce the group_by function

group_by breaks down a dataset into specified groups of rows. When you then apply the verbs above on the resulting object they'll be automatically applied "by group".

Use in conjunction with mutate (to add existing rows to your data frame) or summarise (to create a new data frame)

- `mean`: mean within groups
- `sum`: sum within groups
- `sd`: standard deviation within groups
- `max`: max within groups
- `n()`: number in each group
- `first`: first in group
- `last`: last in group
- `nth(n = 3)`: nth in group (3rd here)
- `tally`: count number in group

OPERATING ON GROUPS I

Which ministry has the largest expenses?

```
df.expense = df %>%
  group_by(paragraf) %>%
  summarise(sum.exp = sum(udgift, na.rm = TRUE)) %>
  arrange(-sum.exp)
```

paragraf	sum.exp
Social- og Indenrigsministeriet	1231214.6
Beskæftigelsesministeriet	1146997.8
Uddannelses- og Forskningsministeriet	296539.2
Min. for Børn, Undervisning og Ligestilling	180955.1
Pensionsvæsenet	139058.0

OPERATING ON GROUPS II

Add `sum.exp` to existing data frame

```
df.2 = df %>%
  group_by(paragraf) %>%
  mutate(sum.exp = sum(udgift, na.rm = TRUE)) %>%
  select(paragraf, udgift, sum.exp)
```

paragraf	udgift	sum.exp
Dronningen	77.7	474.7
Medlemmer af det kongelige hus m.fl.	26.4	161.2
Folketinget	387.8	6137.6
Folketinget	264.2	6137.6
Folketinget	6.5	6137.6

OPERATING ON GROUPS III

You can group by several variables

```
df.expense.2 = df %>%
  group_by(paragraf, aar) %>%
  summarise(sum.exp = sum(udgift, na.rm = TRUE)) %>%
  arrange(sum.exp)
```

| paragraf	aar	sum.exp
Afdrag på statsgælden (netto)	2016	-77832.3
Afdrag på statsgælden (netto)	2015	-32519.9
Afdrag på statsgælden (netto)	2017	0.0
Afdrag på statsgælden (netto)	2018	0.0
Afdrag på statsgælden (netto)	2019	0.0

Let's first calculate yearly expenses at the paragraf level and then calculate mean expenses over the years.

```
df.expense.3 = df %>%
  group_by(paragraf, aar) %>%
  summarise(exp = sum(udgift, na.rm = TRUE)) %>%
  ungroup() %>%
  group_by(paragraf) %>%
  summarise(sum.exp = mean(exp, na.rm = TRUE))
```

paragraf	sum.exp
Afdrag på statsgælden (netto)	-14628.13333
Beholdningsbevægelser mv.	1540.65000
Beskæftigelsesministeriet	191166.30000
Dronningen	79.11667
Energi-, Forsynings- og Klimaministeriet	2433.18333

Look at this dataset

name	alignment	gender	publisher
Magneto	bad	male	Marvel
Storm	good	female	Marvel
Mystique	bad	female	Marvel
Batman	good	male	DC
Joker	bad	male	DC
Catwoman	bad	female	DC
Hellboy	good	male	Dark Horse Comics

And this

publisher	yr_founded
DC	1934
Marvel	1939
Image	1992

```
ijsp = inner_join(superheroes, publishers)
```

name	alignment	gender	publisher	yr_founded
Magneto	bad	male	Marvel	1939
Storm	good	female	Marvel	1939
Mystique	bad	female	Marvel	1939
Batman	good	male	DC	1934
Joker	bad	male	DC	1934
Catwoman	bad	female	DC	1934

```
ljsp = left_join(superheroes, publishers)
```

name	alignment	gender	publisher	yr_founded
Magneto	bad	male	Marvel	1939
Storm	good	female	Marvel	1939
Mystique	bad	female	Marvel	1939
Batman	good	male	DC	1934
Joker	bad	male	DC	1934
Catwoman	bad	female	DC	1934
Hellboy	good	male	Dark Horse Comics	NA

```
superheroes = superheroes %>%
  mutate(seblikes = (publisher == "Marvel"))
publishers = publishers %>%
  mutate(seb = (publisher == "Marvel"))
ij2 = inner_join(superheroes,publishers)

## Joining by: "publisher"
```

name	alignment	gender	publisher	seblikes	yr_founded	seb
Magneto	bad	male	Marvel	TRUE	1939	TRU
Storm	good	female	Marvel	TRUE	1939	TRU
Mystique	bad	female	Marvel	TRUE	1939	TRU
Batman	good	male	DC	FALSE	1934	FAL
Joker	bad	male	DC	FALSE	1934	FAL
Catwoman	bad	female	DC	FALSE	1934	FAL

```
ij2 = inner_join(superheroes, publishers,
                 by=c("publisher"="publisher",
                      "seblikes"="seb"))
```

name	alignment	gender	publisher	seblikes	yr_founded
Magneto	bad	male	Marvel	TRUE	1939
Storm	good	female	Marvel	TRUE	1939
Mystique	bad	female	Marvel	TRUE	1939
Batman	good	male	DC	FALSE	1934
Joker	bad	male	DC	FALSE	1934
Catwoman	bad	female	DC	FALSE	1934

```
fj = superheroes %>%
  full_join(publishers)
```

name	alignment	gender	publisher	yr_founded
Magneto	bad	male	Marvel	1939
Storm	good	female	Marvel	1939
Mystique	bad	female	Marvel	1939
Batman	good	male	DC	1934
Joker	bad	male	DC	1934
Catwoman	bad	female	DC	1934
Hellboy	good	male	Dark Horse Comics	NA
NA	NA	NA	Image	1992

Functions and Iteration

Perhaps most important skill for being effective when working with data: write functions.

Advantages

1. You drastically reduce risk of making mistakes
2. When something exogenous changes, you only need to update code in one place
3. You can give your function an intuitive name that makes your code easier to read

You should write functions to increase your productivity.

```
my_function = function(input1, input2, ..., inputN)
  {
  # define 'output' using input1,...,inputN
  return(output)
  }
```

```
add_numbers = function(x, y){
  z = x + y
  return(z)
}

add_numbers(2, 4)

## [1] 6

add_numbers(2, 6)

## [1] 8
```

Now what

```
add_numbers(2, "y")

## Error in x + y: non-numeric argument to binary operat
```

```
add_numbers = function(x, y){
  if ( !is.numeric(x) || !is.numeric(y)) {
    warning("either 'x' or 'y' is not numeric")
    return(NA)
  }
  else {
    z = x + y
    return(z)
  }
}
```

```
add_numbers(2, 4)

## [1] 6

add_numbers(2, "y")

## Warning in add_numbers(2, "y"): either 'x' or 'y' is

## [1] NA
```

One important skill for being an effective data analyst was being able to write functions. A second is iteration.

Iteration helps you when you need to do the same thing to multiple inputs. For example, repeating the same function on lots of inputs.

Two iteration paradigms

1. Imperative programming (`for` loops, etc.)
2. Functional programming

Imagine that we have this data frame, called **df**

a	b	c	d
0.1801820	-0.8482089	0.7269846	0.0675439
-0.8834473	-0.6590159	0.1412176	0.1648077
-2.4664103	-0.5048838	0.3130126	0.3166068
-0.7819241	-0.9421274	0.8909220	3.6194093
0.4068995	0.3060998	0.2670147	0.2004793

And assume that we want to compute the mean of each column

We could iterate through each column, compute the mean and output the results

```
output = vector()
for (i in 1:ncol(df)){
  output[[i]] = mean(df[, i], na.rm = TRUE)
}
output
```

```
## [1] -0.7089400 -0.5296272  0.4678303  0.8737694
```

We will do Exercise 4.5.1 (*Prediction Based on Betting Markets*) in Imai (2016) using `for` loops.

READ DATA

```
intrade

library("readr")
gh.link = "https://raw.githubusercontent.com/"
user.repo = "kosukeimai/qss/"
branch = "master/"
link = "PREDICTION/intrade08.csv"
data.link = paste0(gh.link, user.repo, branch, link)
intrade.08 = read_csv(data.link)

link = "PREDICTION/intrade12.csv"
data.link = paste0(gh.link, user.repo, branch, link)
intrade.12 = read_csv(data.link)
```

Name	Description
day	Date of the session
statename	Full name of each state
state	Abbreviation of each state
PriceD	Predicted vote share of D Nominee's market
PriceR	Predicted vote share of R Nominee's market
VolumeD	Total session trades of D Nominee's market
VolumeR	Total session trades of R Nominee's market

```
pres

link = "PREDICTION/pres08.csv"
data.link = paste0(gh.link, user.repo, branch, link)
pres.08 = read_csv(data.link)

link = "PREDICTION/pres12.csv"
data.link = paste0(gh.link, user.repo, branch, link)
pres.12 = read_csv(data.link)
```

Name	Description
state.name	Full name of state (only in `pres2008`)
state	Two letter state abbreviation
Obama	Vote percentage for Obama
McCain	Vote percentage for McCain
EV	Number of electoral college votes for this state

```
polls

link = "PREDICTION/polls08.csv"
data.link = paste0(gh.link, user.repo, branch, link)
polls.08 = read_csv(data.link)

link = "PREDICTION/polls12.csv"
data.link = paste0(gh.link, user.repo, branch, link)
polls.12 = read_csv(data.link)
```

Name	Description
state	Abbreviated name of state in which poll was conducted
Obama	Predicted support for Obama (percentage)
Romney	Predicted support for Romney (percentage)
Pollster	Name of organization conducting poll
middate	Middle of the period when poll was conducted

for loops emphasize the objects instead of the functions

```
output = vector()
for (i in 1:ncol(df)){
  output[[i]] = median(df[, i], na.rm = TRUE)
}
```

```
output = vector()
for (i in 1:ncol(df)){
  output[[i]] = mean(df[, i], na.rm = TRUE)
}
```

```r
library("purrr")
map_dbl(df, mean)
```

```
##          a          b          c          d
## -0.7089400 -0.5296272  0.4678303  0.8737694
```

```r
map_dbl(df, median)
```

```
##          a          b          c          d
## -0.7819241 -0.6590159  0.3130126  0.2004793
```

USING YOU OWN FUNCTION

```r
my_function = function(df){
  my.mean = mean(df, na.rm = TRUE)
  my.string = paste("mean is",
                    round(my.mean, 2),
                    sep = ":")
  return(my.string)
}


map_chr(df, my_function)

##                 a                 b                 c
## "mean is:-0.71" "mean is:-0.53"   "mean is:0.47"   "mea
```

```r
library("dplyr")
library("purrr")
library("ggplot2")
```

```
df.08 = intrade.08 %>%
  left_join(pres.08) %>%
  select(-1) %>%
  mutate(
    days.to.election = as.Date("2008-11-04") -
      as.Date(day),
    obama.intmarg = PriceD - PriceR,
    obama.actmarg = Obama - McCain
  ) %>%
  mutate( # gen sign indicator
    obama.intmarg.sign = sign(obama.intmarg),
    obama.actmarg.sign = sign(obama.actmarg)
  )
```

Number of misclassified states in 2008

```
df.08 %>%
  filter(days.to.election == 1) %>%
  filter(obama.intmarg.sign !=
           obama.actmarg.sign) %>%
  count()
```

```
## Source: local data frame [1 x 1]
##
##       n
##   (int)
## 1     2
```

```
df.08.1 = df.08 %>%
  filter(days.to.election > 0 &
           days.to.election <= 90) %>%
  filter(obama.intmarg > 0) %>%
  group_by(days.to.election) %>%
  summarise(
    EV = sum(EV, na.rm = TRUE)
  )
```

```r
library("RcppRoll") # for 'roll_mean' function
df.08.2 = df.08 %>%
  filter(days.to.election > 0 &
         days.to.election <= 90) %>%
  arrange(state, days.to.election) %>%
  group_by(state) %>%
  mutate(
    roll.mean = roll_mean(x = obama.intmarg, 7,
                          align = "right",
                          fill = NA)) %>%
  ungroup() %>%
  filter(roll.mean > 0) %>%
  group_by(days.to.election) %>%
  summarise(
    EV = sum(EV, na.rm = TRUE))
```

String Processing

A regular expression is a pattern that describes a specific set of strings with a common structure. It is heavily used for string matching / replacing in all programming languages, although specific syntax may differ a bit.

Regular expressions typically specify characters (or character classes) to seek out, possibly with information about repeats and location within the string. This is accomplished with the help of metacharacters that have specific meaning: $ * + . ? [ ] ^ { } | ( ) \

They can be difficult

Quantifiers specify how many repetitions of the pattern

- `*`: matches at least 0 times
- `+`: matches at least 1 times
- `?`: matches at most 1 times
- `{n}`: matches exactly *n* times
- `{n,}`: matches at least *n* times
- `{n,m}`: matches between *n* and *m* times

EXAMPLE

```
strings = c("a", "ab", "acb", "accb",
            "acccb", "accccb")
grep("ac*b", strings, value = TRUE)

## [1] "ab"     "acb"    "accb"    "acccb"  "accccb"

grep("ac+b", strings, value = TRUE)

## [1] "acb"    "accb"    "acccb"  "accccb"

grep("ac?b", strings, value = TRUE)

## [1] "ab"  "acb"
```

```r
grep("ac{2}b", strings, value = TRUE)
```

```
## [1] "accb"
```

```r
grep("ac{2,}b", strings, value = TRUE)
```

```
## [1] "accb"    "acccb"  "accccb"
```

```r
grep("ac{2,3}b", strings, value = TRUE)
```

```
## [1] "accb"   "acccb"
```

- `.`: matches any single character, as shown in the first example.
- `[...]`: a character list, matches any one of the characters inside the square brackets. We can also use - inside the brackets to specify a range of characters.
- `[^...]`: an inverted character list, similar to […], but matches any characters except those inside the square brackets.
- `|`: an "or" operator, matches patterns on either side of the |.
- `\`: suppress the special meaning of metacharacters in regular expression
- `^`: matches the start of the string.

```
strings = c("^ab", "ab", "abc", "abd", "abe", "ab 12")
grep("ab.", strings, value = TRUE)

## [1] "abc"    "abd"    "abe"    "ab 12"

grep("ab[c-e]", strings, value = TRUE)

## [1] "abc" "abd" "abe"

grep("ab[^c]", strings, value = TRUE)

## [1] "abd"    "abe"    "ab 12"
```

```r
grep("^ab", strings, value = TRUE)
```

```
## [1] "ab"    "abc"    "abd"    "abe"    "ab 12"
```

```r
grep("\\^ab", strings, value = TRUE)
```

```
## [1] "^ab"
```

```r
grep("abc|abd", strings, value = TRUE)
```

```
## [1] "abc" "abd"
```

CHARACTER CLASSES

Character classes allows to specify entire classes of characters, such as numbers, letters, etc.

- `[:digit:]` or `\d`: digits, 0 1 2 3 4 5 6 7 8 9, equivalent to `[0-9]`.
- `\D`: non-digits, equivalent to `[^0-9]`.
- `[:lower:]`: lower-case letters, equivalent to `[a-z]`.
- `[:upper:]`: upper-case letters, equivalent to `[A-Z]`.
- `[:alpha:]`: alphabetic characters, equivalent to `[[:lower:][:upper:]]` or `[A-z]`.
- `[:alnum:]`: alphanumeric characters, equivalent to `[[:alpha:][:digit:]]` or `[A-z0-9]`.
- `\w`: word characters, equivalent to `[[:alnum:]_]` or `[A-z0-9_]`.
- `\W`: not word, equivalent to `[^A-z0-9_]`.

- `[:blank:]`: blank characters, i.e. space and tab.
- `[:space:]`: space characters: tab, newline, vertical tab, form feed, carriage return, space.
- `\s`: space, ' '.
- `\S`: not space.
- `[:punct:]`: punctuation characters, ! " # $ % & ' ( ) * + , - . / : ; < = > ? @ ^ _ ` { | } ~.

The most consistent way of doing string manipulation in R is by using the stringr package.

```
library("stringr")
```

All functions in stringr start with str_.

All functions are pipeable!

- `str_length`: length of string
- `str_c`: combine string
- `str_sub`: subset string
- `str_to_lower`: string to lower cases

MATCHING WITH REGULAR EXPRESSIONS

Detect matches

```
x = c("apple", "banana", "pear", "213")
x %>% str_detect("e")

## [1]  TRUE FALSE  TRUE FALSE

x %>% str_detect("[0-9]")

## [1] FALSE FALSE FALSE  TRUE

x %>% str_detect("e.r")

## [1] FALSE FALSE  TRUE FALSE
```

Extract matches

```
x %>% str_extract("apple|[0-9]*")
```

```
## [1] "apple" ""        ""        "213"
```

```
x %>% str_extract("a.*")
```

```
## [1] "apple" "anana" "ar"     NA
```

If the string contains more than more element

```
xx = c("apple 123", "banana", "pear", "213")
xx %>% str_extract("apple|[0-9]*")

## [1] "apple"  ""      ""      "213"

xx %>% str_extract("a.*")

## [1] "apple 123" "anana"     "ar"        NA
```

```
xx %>% str_extract_all("apple|[0-9].")

## [[1]]
## [1] "apple" "12"
##
## [[2]]
## character(0)
##
## [[3]]
## character(0)
##
## [[4]]
## [1] "21"
```

Replacing

```
xx %>% str_replace("[0-9]", "\\?")

## [1] "apple ?23" "banana"    "pear"       "?13"

xx %>% str_replace_all("[0-9]", "\\?")

## [1] "apple ???" "banana"    "pear"       "???"

xx %>% str_replace_all(c("1" = "one", "2" = "x"))

## [1] "apple onex3" "banana"     "pear"        "xone3"
```

SPLIT VECTOR

Splitting

```
xx %>% str_split("a")
```

```
## [[1]]
## [1] ""          "pple 123"
##
## [[2]]
## [1] "b" "n" "n" ""
##
## [[3]]
## [1] "pe" "r"
##
## [[4]]
## [1] "213"
```

```r
xx %>% str_split("a", n = 2)

## [[1]]
## [1] ""          "pple 123"
##
## [[2]]
## [1] "b"     "nana"
##
## [[3]]
## [1] "pe" "r"
##
## [[4]]
## [1] "213"
```