

SOCIAL DATA SCIENCE

DATA MANIPULATION

Sebastian Barfort

August 03, 2016

University of Copenhagen
Department of Economics

- dplyr
- tidyr
- purrr
- tidytext
- stringr

“Herein lies the dirty secret about most data scientists’ work – it’s more data munging than deep learning. The best minds of my generation are deleting commas from log files, and that makes me sad. A Ph.D. is a terrible thing to waste.”

Source

TECHNOLOGY

For Big-Data Scientists, 'Janitor Work' Is Key Hurdle to Insights

By STEVE LOHR AUG. 17, 2014

Email

Share

Tweet

Save

More

Technology revolutions come in measured, sometimes foot-dragging steps. The lab science and marketing enthusiasm tend to underestimate the bottlenecks to progress that must be overcome with hard work and practical engineering.

The field known as “big data” offers a contemporary case study. The catchphrase stands for the modern abundance of digital data from many sources — the web, sensors, smartphones and corporate databases — that can be mined with clever software for discoveries and insights. Its promise is smarter, data-driven decision-making in every field. That is why data scientist is the economy’s hot new job.

Yet far too much handcrafted work — what data scientists call “data wrangling,” “data



Monica Rogati, Jawbone's vice president for data science, with Brian Wilt, a senior data scientist. Peter DaSilva for The New York Times

Source

Raw data

The original source of the data

Often hard to use directly for data analysis

You should *never* process your original data


Processed data

Data that is ready for analysis

Data manipulation involves going from *raw* to *processed* data.

This can include merging, subsetting, transforming, etc.

All steps that take you from raw to processed data should be scripted



Tidy data

*Happy families are all alike; every unhappy family is
unhappy in its own way*

Leo Tolstoy

Goal of **tidyr**: take your messy data and turn it into a tidy format

tidy data: observations are in the rows, variables are in the columns

country	year	cases	population
Afghanistan	1999	1815	1842071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1272015272
China	2000	210766	1280008583

variables

country	year	cases	population
Afghanistan	1999	1815	1842071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1272015272
China	2000	210766	1280008583

observations

country	year	cases	population
Afghanistan	1999	1815	1842071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1272015272
China	2000	210766	1280008583

values

```
library("readr")  
library("dplyr")  
library("tidyr")
```

```
df = read_csv("https://raw.githubusercontent.com/hadley/  
head(df, 3)
```

```
## Source: local data frame [3 x 11]
```

```
##
```

```
##   religion <$10k $10-20k $20-30k $30-40k $40-50k $50-
```

```
##      <chr> <int>    <int>    <int>    <int>    <int>    <i
```

```
## 1 Agnostic      27        34        60        81        76
```

```
## 2 Atheist       12        27        37        52        35
```

```
## 3 Buddhist      27        21        30        34        33
```

```
## Variables not shown: $100-150k <int>, >150k <int>, Do
```

```
##    <int>.
```

THE `gather` FUNCTION

Objective: Reshaping wide format to long format

To tidy this data, we need to **gather** the non-variable columns into a two-column key-value pair

`args(gather)`

```
## function (data, key, value, ..., na.rm = FALSE, converge = FALSE,
##           factor_key = FALSE)
## NULL
```

Arguments:

- **data**: data frame
- **key**: column name representing new variable
- **value**: column name representing variable values
- **...**: names of columns to gather (or not gather)

```
df %>% gather(income, frequency, -religion)
```

```
## Source: local data frame [180 x 3]
```

```
##
```

```
##           religion income frequency
##           <chr>   <chr>      <int>
## 1      Agnostic  <$10k         27
## 2      Atheist   <$10k         12
## 3      Buddhist  <$10k         27
## 4      Catholic  <$10k        418
## 5      Don't know/refused <$10k         15
## 6      Evangelical Prot <$10k        575
## 7      Hindu     <$10k          1
## 8      Historically Black Prot <$10k       228
## 9      Jehovah's Witness <$10k        20
## 10     Jewish    <$10k        19
```

```
##
```

This

```
df %>% gather(income, frequency, 2:11)
```

returns the same as

```
df %>% gather(income, frequency, -religion)
```

MORE COMPLICATED EXAMPLE

Billboard data

```
df = read_csv("https://raw.githubusercontent.com/hadley/head(df, 3)
```

```
## Source: local data frame [3 x 81]
```

```
##
```

```
##   year      artist      track  time da
```

```
##   <int>      <chr>      <chr> <chr>
```

```
## 1  2000      2 Pac Baby Don't Cry (Keep...  4:22
```

```
## 2  2000      2Ge+her The Hardest Part Of ...  3:15
```

```
## 3  2000 3 Doors Down      Kryptonite  3:53
```

```
## Variables not shown: wk2 <int>, wk3 <int>, wk4 <int>,
```

```
##   <int>, wk7 <int>, wk8 <int>, wk9 <int>, wk10 <int>,
```

```
##   <int>, wk13 <int>, wk14 <int>, wk15 <int>, wk16 <in
```

```
##   <int>, wk19 <int>, wk20 <int>, wk21 <int>, wk22 <in
```

```
##   <int>, wk25 <int>, wk26 <int>, wk27 <int>, wk2814<in
```

```
##   <int>, wk31 <int>, wk32 <int>, wk33 <int>, wk34 <in
```

TIDYING THE BILLBOARD DATA

To tidy this dataset, we first gather together all the `wk` columns. The column names give the week and the values are the ranks:

```
billboard2 = df %>%  
  gather(week, rank, wk1:wk76, na.rm = TRUE)  
head(billboard2, 3)
```

```
## Source: local data frame [3 x 7]
```

```
##
```

```
##   year      artist      track  time da
```

```
##   <int>      <chr>      <chr> <chr>
```

```
## 1  2000      2 Pac Baby Don't Cry (Keep...  4:22
```

```
## 2  2000      2Ge+her The Hardest Part Of ...  3:15
```

```
## 3  2000 3 Doors Down      Kryptonite  3:53
```

```
## Variables not shown: rank <chr>.
```

What more would we want to do to the data?

DATA CLEANING

Let's turn the week into a numeric variable and create a proper date column

```
billboard3 = billboard2 %>%  
  mutate(  
    week = extract_numeric(week),  
    date = as.Date(date.entered) + 7 * (week - 1)) %>%  
  select(-date.entered) %>%  
  arrange(artist, track, week)  
head(billboard3, 3)
```

```
## Source: local data frame [3 x 7]
```

```
##
```

```
##   year artist          track  time  week  r  
##   <int> <chr>          <chr> <chr> <dbl> <c  
## 1  2000    2 Pac Baby Don't Cry (Keep... 4:22    1  
## 2  2000    2 Pac Baby Don't Cry (Keep... 4:22    2/61  
## 3  2000    2 Pac Baby Don't Cry (Keep... 4:22    3
```


EVEN MORE COMPLICATED EXAMPLE

After gathering columns, the key column is sometimes a combination of multiple underlying variable names.

```
df = read_csv("https://raw.githubusercontent.com/hadley/head(df, 3)
```

```
## Source: local data frame [3 x 22]
```

```
##
```

```
##      iso2  year   m04   m514   m014 m1524 m2534 m3544 m45
```

```
##      <chr> <int> <int> <int> <int> <int> <int> <int> <in
```

```
## 1      AD  1989     NA     NA     NA     NA     NA     NA
```

```
## 2      AD  1990     NA     NA     NA     NA     NA     NA
```

```
## 3      AD  1991     NA     NA     NA     NA     NA     NA
```

```
## Variables not shown: f04 <int>, f514 <int>, f014 <int>
```

```
##      <int>, f3544 <int>, f4554 <int>, f5564 <int>, f65 <
```

Question: what are the variables here?

The dataset comes from the World Health Organisation, and records the counts of confirmed tuberculosis cases by country, year, and demographic group. The demographic groups are broken down by sex (m, f) and age (0-14, 15-25, 25-34, 35-44, 45-54, 55-64, unknown).

GATHERING THE NON-VARIABLE COLUMNS

```
tb2 = df %>%  
  gather(demo, n, -iso2, -year, na.rm = TRUE)  
head(tb2, 3)
```

```
## Source: local data frame [3 x 4]
```

```
##
```

```
##   iso2   year demo      n
```

```
##   <chr> <int> <chr> <int>
```

```
## 1    AD  2005  m04      0
```

```
## 2    AD  2006  m04      0
```

```
## 3    AD  2008  m04      0
```

Is this dataset tidy?

SEPARATING THE demo VARIABLE

`separate` makes it easy to split a compound variables into individual variables. You can either pass it a regular expression to split on or a vector of character positions. In this case we want to split after the first character.

```
tb3 = tb2 %>%  
  separate(demo, c("sex", "age"), 1)  
head(tb3, 3)
```

```
## Source: local data frame [3 x 5]  
##  
##   iso2  year  sex  age    n  
##   <chr> <int> <chr> <chr> <int>  
## 1    AD  2005    m   04     0  
## 2    AD  2006    m   04     0  
## 3    AD  2008    m   04     0
```

Question: Compare `tb3` to the original data frame (`df`). What are

RESHAPING FROM LONG TO WIDE FORMAT

There are times when we are required to turn long formatted data into wide formatted data. The **spread** function spreads a key-value pair across multiple columns.

args(spread)

```
## function (data, key, value, fill = NA, convert = FALSE,  
##       sep = NULL)  
## NULL
```

- **data**: data frame
- **key**: column values to convert to multiple columns
- **value**: single column values to convert to multiple columns' values
- **fill**: If there isn't a value for every combination of the other variables and the key column, this value will be substituted

spread IN ACTION

```
tb3.wide = tb3 %>% spread(age, n)
tb3.wide
```

```
## Source: local data frame [4,885 x 13]
```

```
##
```

##	iso2	year	sex	014	04	1524	2534	3544	4
##	<chr>	<int>	<chr>	<int>	<int>	<int>	<int>	<int>	<int>
## 1	AD	1996	f	0	NA	1	1	0	
## 2	AD	1996	m	0	NA	0	0	4	
## 3	AD	1997	f	0	NA	1	2	3	
## 4	AD	1997	m	0	NA	0	1	2	
## 5	AD	1998	m	0	NA	0	0	1	
## 6	AD	1999	f	0	NA	0	0	1	
## 7	AD	1999	m	0	NA	0	0	1	
## 8	AD	2000	m	0	NA	0	1	0	
## 9	AD	2001	m	0	NA	NA	NA	2	
## 10	AD	2002	f	0	NA	1	0	0	

In this part of the lecture we will work with the Danish federal budget proposal for 2016

```
library("readr")  
df = read_csv("https://raw.githubusercontent.com/sebasti
```

Some nice guy has already cleaned this data for you

Useful functions:

- `str`: displays the structure of your data frame
- `head`: displays the first rows
- `summary`: gives summary statistics
- `glimpse` (from the `dplyr` package): modern alternative to `str`


```
str(df)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    8430 obs  
## $ paragraf      : chr  "Dronningen" "Medlemmer af det k  
## $ hovedomrode: chr  "Statsydelse" "Årpenge mv." "Udg  
## $ aktivitet     : chr  "Statsydelse" "Årpenge mv." "Fol  
## $ hovedkonto   : chr  "Statsydelse" "Årpenge mv." "Fol  
## $ aar          : int   2014 2014 2014 2014 2014 2014 20  
## $ udgift       : num   77.7 26.4 387.8 264.2 6.5 ...
```

head

```
head(df)
```

```
## Source: local data frame [6 x 6]
```

```
##
```

```
##
```

```
paragraf
```

```
h
```

```
##
```

```
<chr>
```

```
## 1
```

```
Dronningen
```

```
S
```

```
## 2 Medlemmer af det kongelige hus m.fl.
```

```
Å
```

```
## 3
```

```
Folketinget Udgifter ved F
```

```
## 4
```

```
Folketinget Udgifter ved F
```

```
## 5
```

```
Folketinget Udgifter ved F
```

```
## 6
```

```
Folketinget Udgifter ved F
```

```
## Variables not shown: aktivitet <chr>, hovedkonto <chr>
```

```
## <dbl>.
```

summary

```
summary(df)
```

```
##      paragraf      hovedomrode      aktivitet
## Length:8430      Length:8430      Length:8430
## Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character
##
##
##      hovedkonto      aar      udgift
## Length:8430      Min.   :2014      Min.   : -306424.10
## Class :character 1st Qu.:2015      1st Qu.:      0.00
## Mode  :character Median :2016      Median :      7.00
##                  Mean  :2016      Mean  :    223.18
##                  3rd Qu.:2018      3rd Qu.:    89.95
##                  Max.   :2019      Max.   : 132541.40
```

```
library("dplyr")  
glimpse(df)
```

```
## Observations: 8,430  
## Variables: 6  
## $ paragraf      (chr) "Dronningen", "Medlemmer af det k  
## $ hovedomrode   (chr) "Statsydelse", "Årpenge mv.", "Ud  
## $ aktivitet     (chr) "Statsydelse", "Årpenge mv.", "Fo  
## $ hovedkonto    (chr) "Statsydelse", "Årpenge mv.", "Fo  
## $ aar           (int) 2014, 2014, 2014, 2014, 2014, 201  
## $ udgift        (dbl) 77.7, 26.4, 387.8, 264.2, 6.5, 2.
```

Data Manipulation (`dplyr`)

Many data analysis problems involve the application of a *split-apply-combine strategy*, where you break up a big problem into manageable pieces, operate on each piece independently and then put the pieces back together



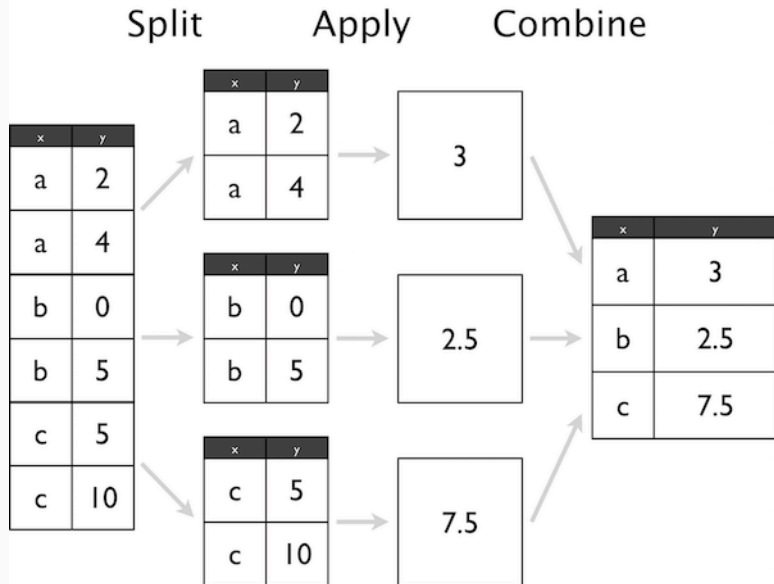
Journal of Statistical Software

April 2011, Volume 40, Issue 1.

<http://www.jstatsoft.org/>

The Split-Apply-Combine Strategy for Data Analysis

Hadley Wickham
Rice University



dplyr: (efficiently) split-apply-combine for data frames

Verbs a verb is a function that takes a data frame as it's first argument

- **filter**: select rows
- **arrange**: order rows
- **select**: select columns
- **rename**: rename columns
- **distinct**: find distinct rows
- **mutate**: add new variables
- **summarise**: summarize across a data set
- **sample_n**: sample from a data set

THE filter FUNCTION I

```
filter(df, udgift == min(udgift))
```

```
## Source: local data frame [1 x 6]
```

```
##
```

```
##           paragraf           hovedomrode  
##           <chr>           <chr>
```

```
## 1 Skatter og afgifter Skatter på indkomst og formue P
```

```
## Variables not shown: hovedkonto <chr>, aar <int>, udg
```

```
filter(df, paragraf == "Skatter og afgifter ")
```


```
## Source: local data frame [0 x 6]
```

```
##
```

```
## Variables not shown: paragraf <chr>, hovedomrode <chr>
```

```
##   hovedkonto <chr>, aar <int>, udgift <dbl>.
```

LOGICAL OPERATORS



COMBINING CONDITIONS: and

You can easily combine conditions

```
filter(df, paragraf == "Skatter og afgifter " & aktivite
```

```
## Source: local data frame [0 x 6]
```

```
##
```

```
## Variables not shown: paragraf <chr>, hovedomrode <chr>
```

```
##   hovedkonto <chr>, aar <int>, udgift <dbl>.
```

```
filter(df, paragraf == "Skatter og afgifter " | aktivite  
  
## Source: local data frame [0 x 6]  
##  
## Variables not shown: paragraf <chr>, hovedomrode <chr>  
##   hovedkonto <chr>, aar <int>, udgift <dbl>.
```

THE select FUNCTION

```
select(df, aar, udgift)
```

```
## Source: local data frame [8,430 x 2]
```

```
##
```

```
##      aar udgift
```

```
##    <int> <dbl>
```

```
## 1    2014   77.7
```

```
## 2    2014   26.4
```

```
## 3    2014  387.8
```

```
## 4    2014  264.2
```

```
## 5    2014    6.5
```

```
## 6    2014    2.2
```

```
## 7    2014   63.3
```

```
## 8    2014    5.0
```

```
## 9    2014  202.5
```

```
## 10   2014   76.7
```

```
## .. ... ..
```

THE arrange FUNCTION

```
arrange(df, hovedomrode, udgift)
```

```
## Source: local data frame [8,430 x 6]
```

```
##
```

```
##                                     paragraf  
##                                     <chr>
```

```
## 1  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 2  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 3  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 4  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 5  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 6  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 7  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 8  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 9  Min. for Børn, Undervisning og Ligestilling Admini
```

```
## 10 Min. for Børn, Undervisning og Ligestilling Admini
```

```
## ..
```

ARRANGE BY NUMERIC VARIABLE

```
arrange(df, -aar)
```

```
## Source: local data frame [8,430 x 6]
```

```
##
```

```
##           paragraf
```

```
##           <chr>
```

```
## 1           Dronningen
```

```
## 2 Medlemmer af det kongelige hus m.fl.
```

```
## 3           Folketinget          Udg
```

```
## 4           Folketinget          Udg
```

```
## 5           Folketinget          Udg
```

```
## 6           Folketinget          Udg
```

```
## 7           Folketinget          Fo
```

```
## 8           Folketinget Statsrevisore
```

```
## 9           Folketinget Statsrevisore
```

```
## 10          Folketinget          F
```

```
## ..
```

```
...
```

THE mutate FUNCTION

mutate let's you add new variables to your data frame

```
df.mutated = mutate(df, newVar = udgift/2)  
select(df.mutated, newVar, udgift)
```

```
## Source: local data frame [8,430 x 2]
```

```
##
```

```
##      newVar udgift
```

```
##      <dbl> <dbl>
```

```
## 1    38.85   77.7
```

```
## 2    13.20   26.4
```

```
## 3   193.90  387.8
```

```
## 4   132.10  264.2
```

```
## 5     3.25    6.5
```

```
## 6     1.10    2.2
```

```
## 7    31.65   63.3
```

```
## 8     2.50    5.0
```

```
## 9    101.25  202.5
```


THE sample_n FUNCTION

```
sample_n(df, 3)
```

```
## Source: local data frame [3 x 6]
```

```
##
```

```
##
```

```
paragraf
```

```
hov
```

```
##
```

```
<chr>
```

```
## 1
```

```
Sundheds- og Ældreministeriet
```

```
Syge
```

```
## 2
```

```
Transport- og Bygningsministeriet Trafik og byg
```

```
## 3
```

```
Uddannelses- og Forskningsministeriet Fælles
```

```
## Variables not shown: aktivitet <chr>, hovedkonto <chr>
```

```
##
```

```
<dbl>.
```

magrittr::

%>%

The pipe operator `%>%` (RStudio has keyboard shortcuts, learn to use them!) let's you write sequences instead of nested functions

```
x %>% f(y) -> f(x, y)
```

```
x %>% f(z, .) -> f(z, x)
```

Read `%>%` as “then”. First do this, *then* do this, etc...

It's implemented in R by a **Danish econometrician**

INTUITION

```
enjoy(cool(bake(shape(beat(append(bowl(rep("flour", 2),  
"yeast", "water", "milk", "oil"), "flour", until =  
"soft"), duration = "3mins"), as = "balls", style =  
"slightly-flat"), degrees = 200, duration = "15mins"),  
duration = "5mins"))
```

```
bowl(rep("flour", 2), "yeast", "water", "milk", "oil") %>%  
  append("flour", until = "soft") %>%  
  beat(duration = "3mins") %>%  
  shape(as = "balls", style = "slightly-flat") %>%  
  bake(degrees = 200, duration = "15mins") %>%  
  cool(buns, duration = "5mins") %>%  
  enjoy()
```

source

COMBINING `dplyr` AND THE PIPE

`dplyr` is designed to work with the pipe.

So

```
df %>%  
  select(aar, udgift) %>%  
  filter(aar == 2014)
```

returns the same as

```
filter(select(df, aar, udgift), aar == 2014)
```

EXAMPLE

Show me a random sample of the data from 2014, where `paragraf == Folketinget` and `udgift` is above the mean.

```
df.1 = filter(df, aar == 2014 & paragraf == "Folketinget")
df.2 = filter(df.1, udgift > mean(udgift, na.rm = TRUE))
df.3 = sample_n(df.2, 3)
df.3
```

```
## Source: local data frame [3 x 6]
##
##      paragraf                                hovedomrode
##      <chr>                                     <chr>
## 1 Folketinget                Udgifter ved Folketinget
## 2 Folketinget Statsrevisorerne og Rigsrevisionen
## 3 Folketinget                Udgifter ved Folketinget
## Variables not shown: aktivitet <chr>, hovedkonto <chr>
##      <dbl>.
```

```
df %>%
  filter(aar == 2014 & paragraf == "Folketinget") %>%
  filter(udgift > mean(udgift, na.rm = TRUE)) %>%
  sample_n(3)

## Source: local data frame [3 x 6]
##
##      paragraf                                hovedomrode
##      <chr>                                <chr>
## 1 Folketinget          Udgifter ved Folketinget
## 2 Folketinget Statsrevisorerne og Rigsrevisionen
## 3 Folketinget          Udgifter ved Folketinget
## Variables not shown: aktivitet <chr>, hovedkonto <chr>
##      <dbl>.
```

Note how readable the code is. Almost like a grammar of data manipulation?

So far, we have primarily learned how to manipulate data frames.

The **dplyr** package becomes really powerful when we introduce the **group_by** function

group_by breaks down a dataset into specified groups of rows. When you then apply the verbs above on the resulting object they'll be automatically applied "by group".

Use in conjunction with **mutate** (to add existing rows to your data frame) or **summarise** (to create a new data frame)

COMMON mutate/summarise OPTIONS

- `mean`: mean within groups
- `sum`: sum within groups
- `sd`: standard deviation within groups
- `max`: max within groups
- `n()`: number in each group
- `first`: first in group
- `last`: last in group
- `nth(n = 3)`: nth in group (3rd here)

group_by IN ACTION I

Which ministry has the largest expenses?

```
df %>% filter(udgift >= 0) %>% group_by(paragraf) %>%  
  summarise(totale.udgifter = sum(udgift, na.rm = TRUE))  
  arrange(-totale.udgifter)
```

```
## Source: local data frame [28 x 2]
```

```
##
```

```
##               paragraf totale  
##               <chr>
```

```
## 1      Beskæftigelsesministeriet
```

```
## 2      Social- og Indenrigsministeriet
```

```
## 3      Uddannelses- og Forskningsministeriet
```

```
## 4      Min. for Børn, Undervisning og Ligestilling
```

```
## 5      Finansministeriet
```

```
## 6      Pensionsvæsenet
```

```
## 7      Forsvarsministeriet
```

group_by IN ACTION II

Add `totale.udgifter` to the existing data frame

```
df %>% filter(udgift >= 0) %>% group_by(paragraf) %>%  
  mutate(totale.udgifter = sum(udgift, na.rm = TRUE)) %>%  
  select(aar, udgift, totale.udgifter)
```

```
## Source: local data frame [7,609 x 4]
```

```
## Groups: paragraf [28]
```

```
##
```

```
##           paragraf    aar udgift  
##           (chr) (int) (dbl)  
## 1      Dronningen  2014   77.7  
## 2 Medlemmer af det kongelige hus m.fl.  2014   26.4  
## 3      Folketinget  2014  387.8  
## 4      Folketinget  2014  264.2  
## 5      Folketinget  2014    6.5  
## 6      Folketinget  2014    2.2
```

group_by IN ACTION III

You can group by several variables

```
df %>% filter(udgift >= 0) %>% group_by(aar, paragraf) %  
  summarise(totale.udgifter = sum(udgift, na.rm = TRUE))  
  arrange(-totale.udgifter)
```

```
## Source: local data frame [168 x 3]
```

```
## Groups: aar [6]
```

```
##
```

```
##      aar                                paragraf  
##      (int)                                (chr)
```

```
## 1    2014                                Beskæftigelsesministeriet
```

```
## 2    2014                                Social- og Indenrigsministeriet
```

```
## 3    2014                                Afdrag på statsgælden (netto)
```

```
## 4    2014                                Uddannelses- og Forskningsministeriet
```

```
## 5    2014 Min. for Børn, Undervisning og Ligestilling
```

```
## 6    2014                                Renter
```

group_by IN ACTION IV

You can group by several variables

```
df %>% filter(udgift >= 0) %>% group_by(paragraf, hovedo  
  summarise(totale.udgifter = sum(udgift, na.rm = TRUE))  
  arrange(-totale.udgifter)
```

```
## Source: local data frame [120 x 3]
```

```
## Groups: paragraf [28]
```

```
##
```

```
##               paragraf
```

```
##               (chr)
```

```
## 1 Afdrag på statsgælden (netto)
```

```
## 2   Beholdningsbevægelser mv.
```

```
## 3   Beholdningsbevægelser mv.
```


```
## 4   Beholdningsbevægelser mv.
```

```
## 5   Beskæftigelsesministeriet
```

```
## 6   Beskæftigelsesministeriet
```

MERGING DATA SETS

Merging two data sets can be tricky and depends on your needs. It's important to think about what you want before joining.



SUPERHERO EXAMPLE

link

```
superheroes = c("      name, alignment, gender,          pu  
    " Magneto,      bad,   male,           Marvel",  
    "   Storm,    good, female,           Marvel",  
    "Mystique,     bad, female,           Marvel",  
    "   Batman,    good,   male,           DC",  
    "   Joker,     bad,   male,           DC",  
    "Catwoman,     bad, female,           DC",  
    " Hellboy,     good,   male, Dark Horse Comics")
```

```
superheroes = read.csv(text = superheroes, strip.white =  
head(superheroes)
```

```
##      name alignment gender publisher  
## 1  Magneto      bad   male    Marvel  
## 2   Storm    good female    Marvel  
## 3 Mystique     bad female    Marvel
```

```
publishers = c("publisher", yr_founded",  
              "      DC,      1934",  
              "    Marvel,    1939",  
              "    Image,    1992")  
publishers = read.csv(text = publishers, strip.white = T)  
head(publishers)
```

```
## publisher yr_founded  
## 1      DC      1934  
## 2   Marvel    1939  
## 3    Image    1992
```


INNER JOIN

```
ijsp = inner_join(superheroes, publishers)
ijsp
```

##	name	alignment	gender	publisher	yr_founded
## 1	Magneto	bad	male	Marvel	1939
## 2	Storm	good	female	Marvel	1939
## 3	Mystique	bad	female	Marvel	1939
## 4	Batman	good	male	DC	1934
## 5	Joker	bad	male	DC	1934
## 6	Catwoman	bad	female	DC	1934

LEFT JOIN

```
ljsp = left_join(superheroes, publishers)
ljsp
```

##	name	alignment	gender	publisher	yr_foun
## 1	Magneto	bad	male	Marvel	1
## 2	Storm	good	female	Marvel	1
## 3	Mystique	bad	female	Marvel	1
## 4	Batman	good	male	DC	1
## 5	Joker	bad	male	DC	1
## 6	Catwoman	bad	female	DC	1
## 7	Hellboy	good	male	Dark Horse Comics	

MERGING DIFFERENT NAMES

```
superheroes = mutate(superheroes,  
                      seblikes = (publisher=="Marvel"))  
publishers = mutate(publishers,  
                    seb = (publisher == "Marvel"))  
ij2 = inner_join(superheroes,publishers)  
ij2
```

```
##      name alignment gender publisher seblikes yr_fou  
## 1  Magneto      bad   male   Marvel     TRUE  
## 2   Storm     good female   Marvel     TRUE  
## 3 Mystique     bad female   Marvel     TRUE  
## 4   Batman     good   male     DC     FALSE  
## 5    Joker     bad   male     DC     FALSE  
## 6 Catwoman     bad female     DC     FALSE
```

MERGING DIFFERENT NAMES

```
ij2 = inner_join(superheroes,publishers,  
                  by=c("publisher"="publisher",  
                       "seblikes"="seb"))
```

```
ij2
```

##		name	alignment	gender	publisher	seblikes	yr_fou
## 1		Magneto	bad	male	Marvel	TRUE	
## 2		Storm	good	female	Marvel	TRUE	
## 3		Mystique	bad	female	Marvel	TRUE	
## 4		Batman	good	male	DC	FALSE	
## 5		Joker	bad	male	DC	FALSE	
## 6		Catwoman	bad	female	DC	FALSE	

Iteration